



(실습) 7차수



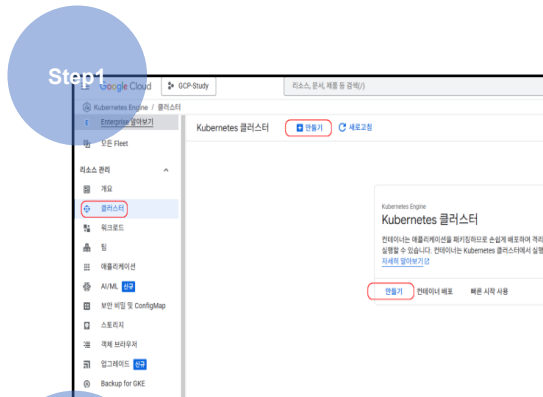
- I GKE
- II GCP Cloud Build
- III CLI 활용 클러스터 생성 및 컨테이너 배포
- IV IaC(Infrastructure as Code)

I GKE - 클러스터 생성

교육 서비스

[GKE Autopilot 모드로 클러스터 만들기]

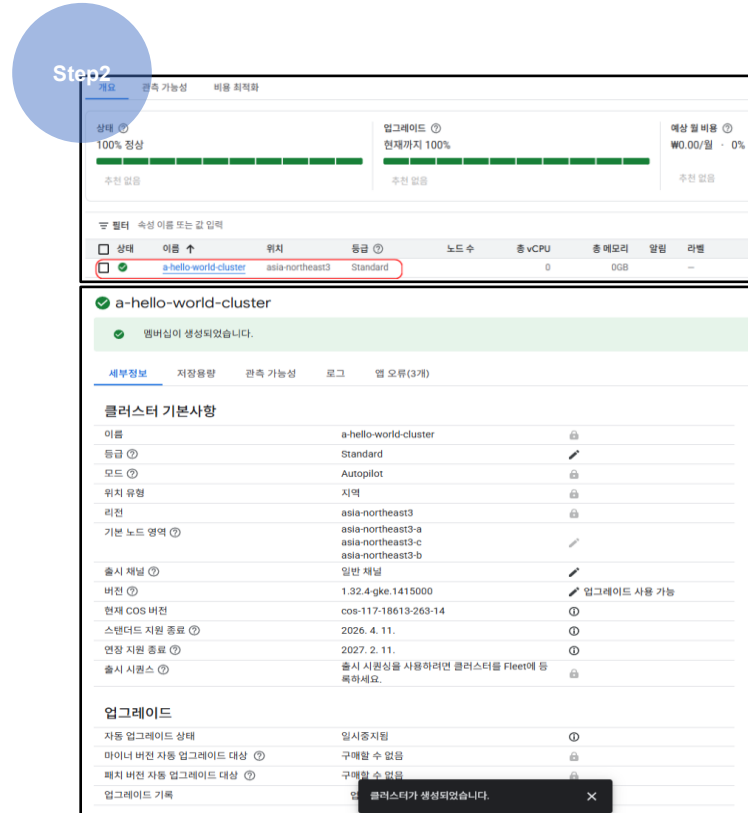
- Autopilot 모드에서 Google은 확장, 보안, 기타 사전 구성된 설정을 포함하여 클러스터 구성을 관리
- Autopilot 모드의 클러스터는 대부분의 프로덕션 워크로드를 실행하고 Kubernetes 매니페스트에 따라 컴퓨팅 리소스를 프로비저닝하도록 최적화되어 있음



[그림 1] 클러스터 만들기



[그림 2] 클러스터 기본사항 설정



[그림 3] 클러스터 생성 결과

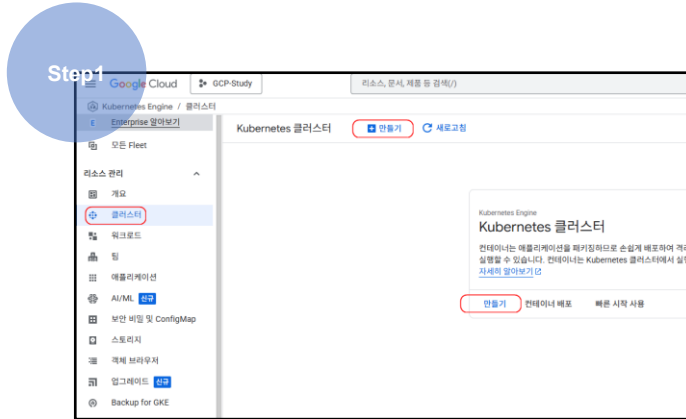
I GKE - 클러스터 생성

교육 서비스

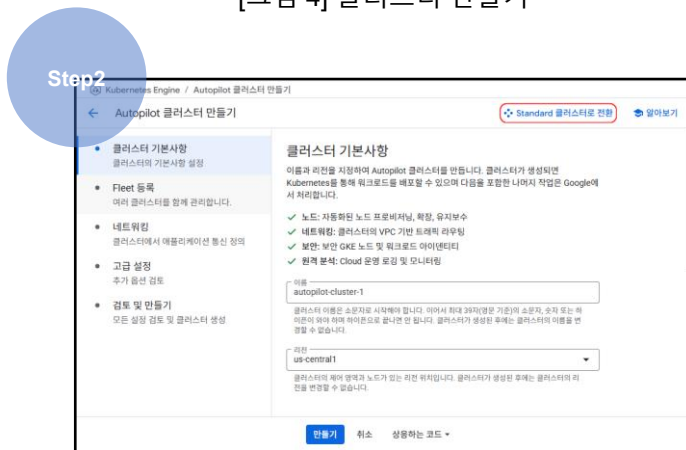


[GKE Standard 모드로 클러스터 만들기]

- Autopilot 모드에 비해 추가 기본설정(node 등) 적용 가능



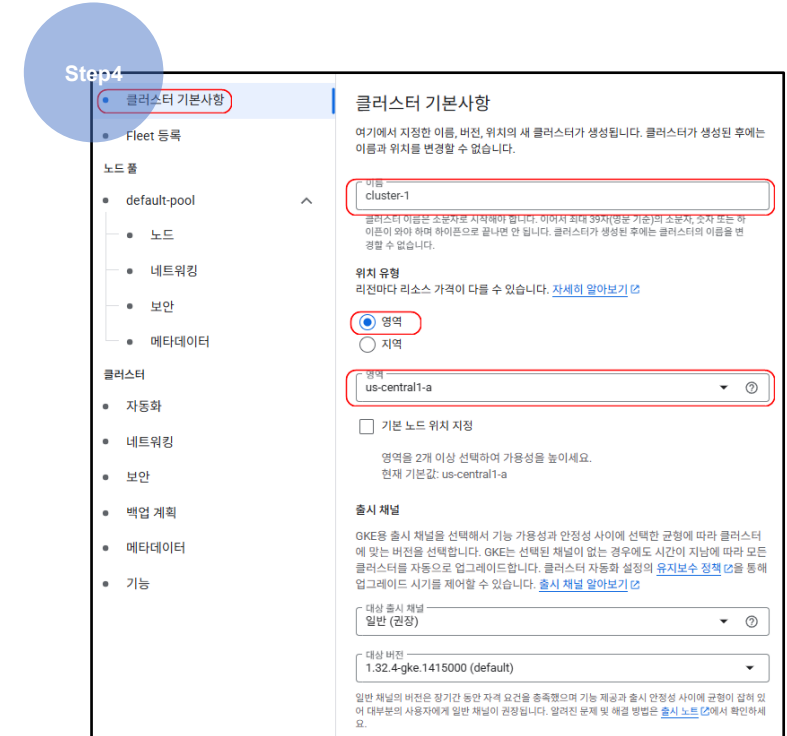
[그림 4] 클러스터 만들기



[그림 5] standard 모드 전환



[그림 6] standard 메뉴



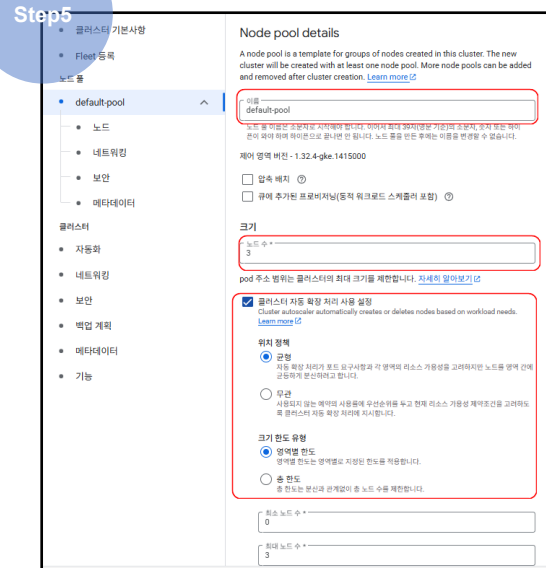
[그림 7] 클러스터 기본사항 설정

I GKE - 클러스터 생성

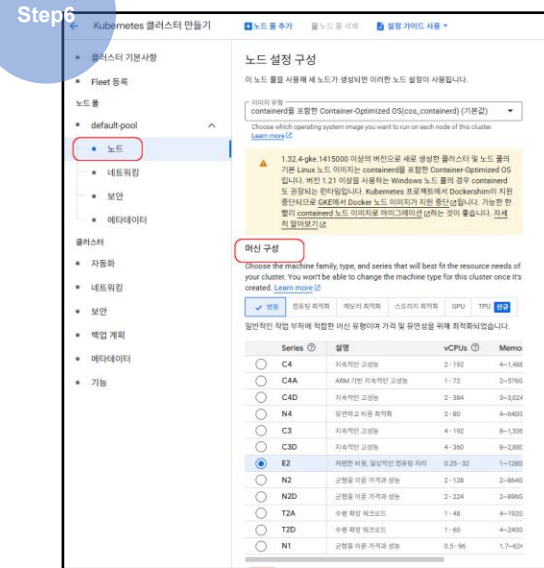
교육 서비스

[GKE standard 모드로 클러스터 만들기]

- Autopilot 모드에 비해 추가 기본설정 적용 가능



[그림 8] default pool 설정




[그림 9] 노드 설정



[그림 10] 머신유형 설정



[그림 11] 클러스터 생성

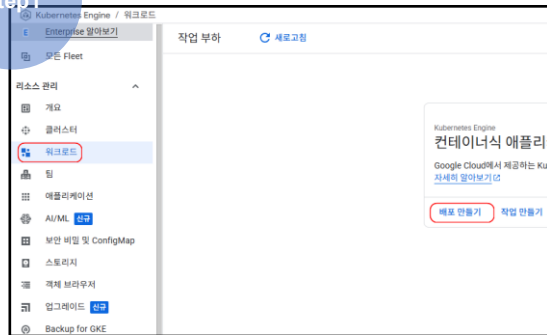


[그림 12] 클러스터 생성 확인

[클러스터에 샘플 앱 배포]

- Google에서 제공되고 Artifact Registry에 컨테이너로 저장되는 샘플 'hello world' 웹 앱을 배포

Step1



[그림 13] 워크로드 배포 만들기

Step2



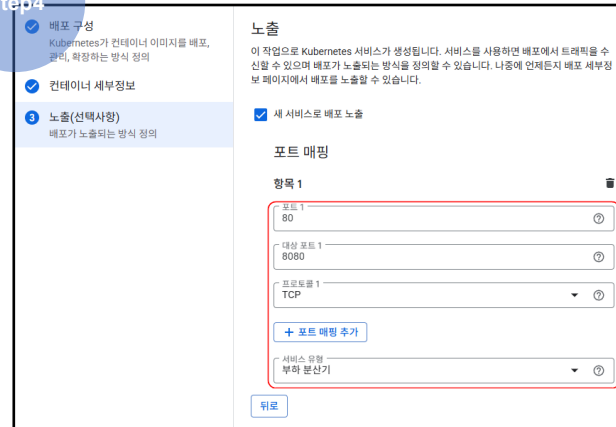
[그림 14] 배포구성 설정

Step3



[그림 15] 컨테이너 세부정보 설정

Step4



[그림 16] 노출 설정



[클러스터에 샘플 앱 배포]

- 샘플 'hello world' 웹앱 yaml 설정 내용

```
---
apiVersion: "apps/v1"
kind: "Deployment"
metadata:
  name: "hello-world-app"
  namespace: "default"
  labels:
    app: "hello-world-app"
spec:
  replicas: 3
  selector:
    matchLabels:
      app: "hello-world-app"
  template:
    metadata:
      labels:
        app: "hello-world-app"
    spec:
      containers:
        - name: "hello-app-1"
          image: "us-docker.pkg.dev/google-samples/containers/gke/hello-app:1.0"
```

[그림 17] hello-world-app-deployment.yml

```
---
apiVersion: "autoscaling/v2"
kind: "HorizontalPodAutoscaler"
metadata:
  name: "hellow-world-app-hpa-dubi"
  namespace: "default"
  labels:
    app: "hello-world-app"
spec:
  scaleTargetRef:
    kind: "Deployment"
    name: "hello-world-app"
    apiVersion: "apps/v1"
  minReplicas: 1
  maxReplicas: 5
  metrics:
    - type: "Resource"
      resource:
        name: "cpu"
        target:
          type: "Utilization"
          averageUtilization: 80
```

[그림 18] hello-world-app-hpa.yml

```
---
apiVersion: "v1"
kind: "Service"
metadata:
  name: "hello-world-app-service"
  namespace: "default"
  labels:
    app: "hello-world-app"
spec:
  ports:
    - protocol: "TCP"
      port: 80
      targetPort: 8080
  selector:
    app: "hello-world-app"
  type: "LoadBalancer"
  loadBalancerIP: ""
```

[그림 19] hello-world-app-svc-lb.yml

[클러스터에 샘플 앱 배포]

- 샘플 'hello world' 웹앱을 배포 완료 상태 확인

Step1

작업 부하 새로고침 배포 작업 만들기 삭제

클러스터 네임스페이스 재설정 Save

개요 관측 가능성 비용 최적화

필터 시스템 객체: False 작업 부하 필터링

<input type="checkbox"/>	이름 ↑	상태	유형	pod	네임스페이스	클러스터
<input checked="" type="checkbox"/>	hello-world-app	OK	Deployment	3/3	default	a-hello-world-cluster

[그림 20] 배포 완료 상태

Step3

← → ↺ ⚠ 주의 요함 34.64.218.53

```

Hello, world!
Version: 1.0.0
Hostname: hello-world-app-6b6cb88cb7-fjww4
    
```

[그림 22] 배포 완료 상태(상세)

Step2

hello-world-app

개요 세부정보 관측 가능성 버전 기록 이벤트 로그 앱 오류(0개) YAML

CPU 메모리 디스크

클러스터 a-hello-world-cluster
네임스페이스 default
앱 hello-world-app app.kubernetes.io/managed-by: cloud-console

로그 Container logs Audit logs
문제본 3개 업데이트됨, 3개 준비됨, 3개 사용 가능, 0개 사용 불가
pod 사양 버전 1, 컨테이너 hello-app-1
수행형 포트 자동 확장 OK
치리
수직형 포트 자동 확장 구성된되지 않음
치리

활성 버전	이름	상태	요약	생성일	실행 중인 pod/총 pod 수
1	hello-world-app-6b6cb88cb7	OK	hello-app-1: us-docker.pkg.dev/google-samples/containers/gke/hello-app:1.0	2025. 7. 4. PM 2:54:46	3/3

관리형 pod	수정 버전	이름	상태	재시작	생성일 ↑
1	hello-world-app-6b6cb88cb7-27mjs	Running	0	2025. 7. 4. PM 2:54:47	
1	hello-world-app-6b6cb88cb7-rw5x2	Running	0	2025. 7. 4. PM 2:54:47	
1	hello-world-app-6b6cb88cb7-fjww4	Running	0	2025. 7. 4. PM 2:54:47	

노출 중인 서비스

이름 ↑	유형	엔드포인트
hello-world-app-service	부하 분산기	34.64.218.53:80

[그림 21] 배포 완료 상태(상세)

I GKE – 노드 관리

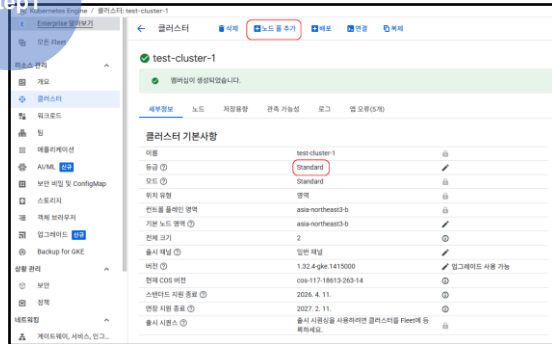
교육 서비스



[노드 풀(Node Pool) 추가]

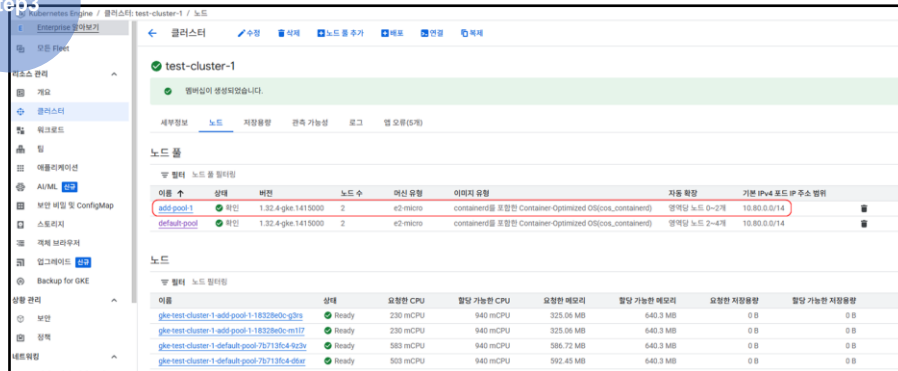
- 클러스터 > 대상 클러스터 이름 선택 > 노드 풀 추가
- 노드 풀(리스트 확인)

Step1



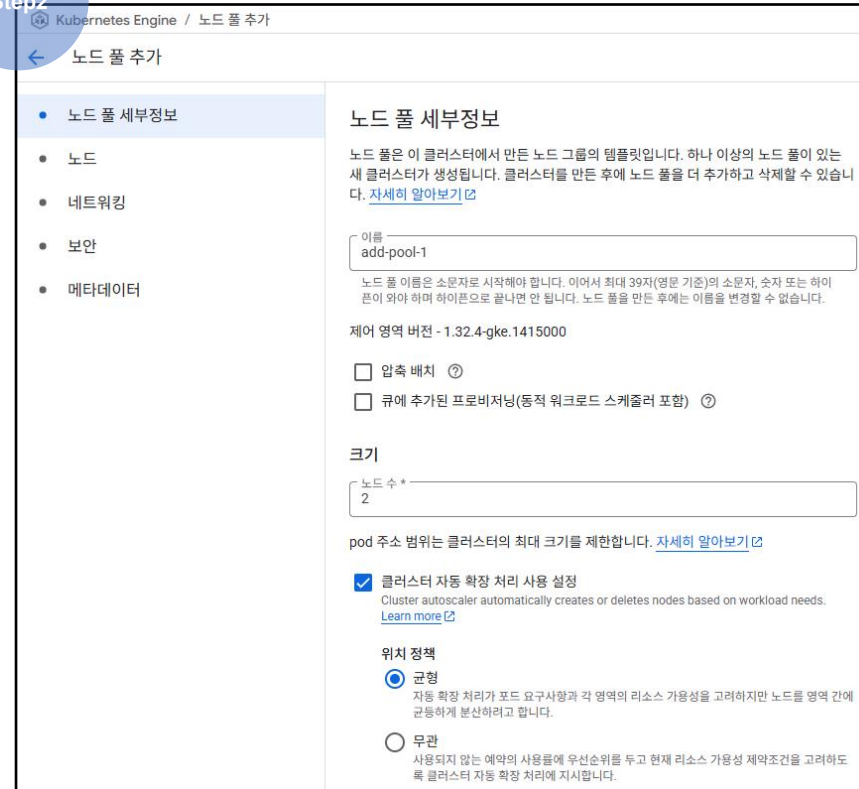
[그림 23] 노드풀 추가

Step3



[그림 25] 노드 풀 추가 확인

Step2

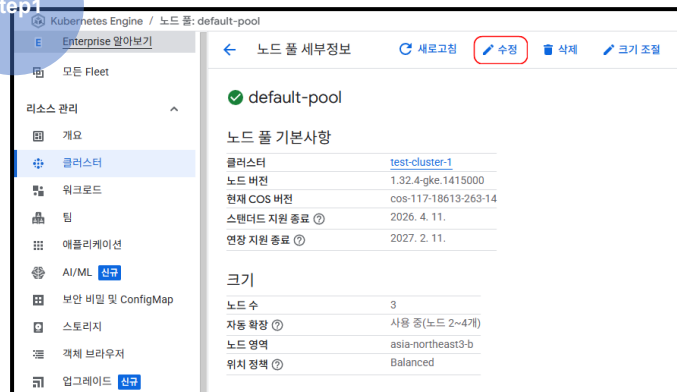


[그림 24] 노드 풀 세부정보 설정

[노드 풀(Node Pool) 수정]

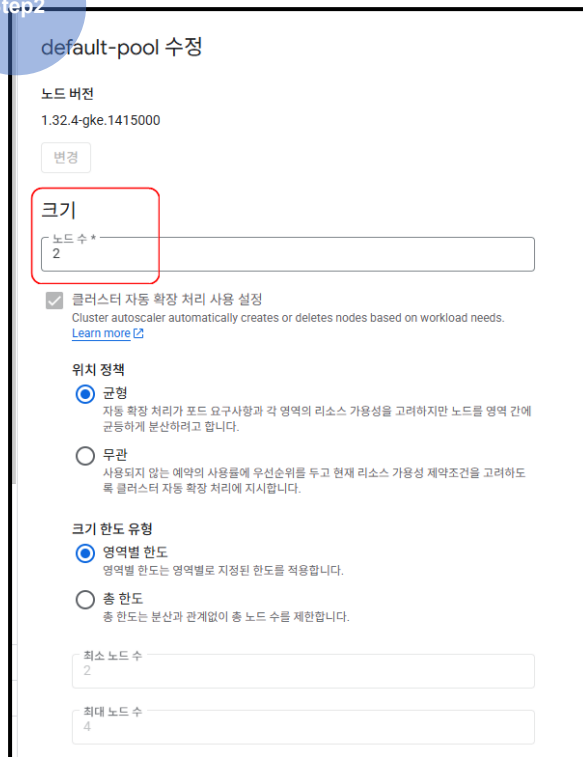
- 클러스터 > 대상 클러스터 이름 선택 > 노드(탭) > 삭제 대상 노드 풀(삭제)
- 노드 풀(리스트 확인)

Step1



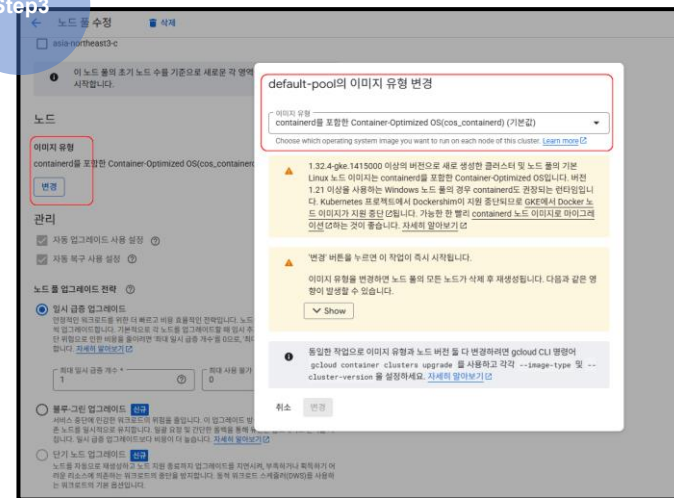
[그림 26] 노드 풀 수정

Step2



[그림 27] 노드 풀 수정

Step3

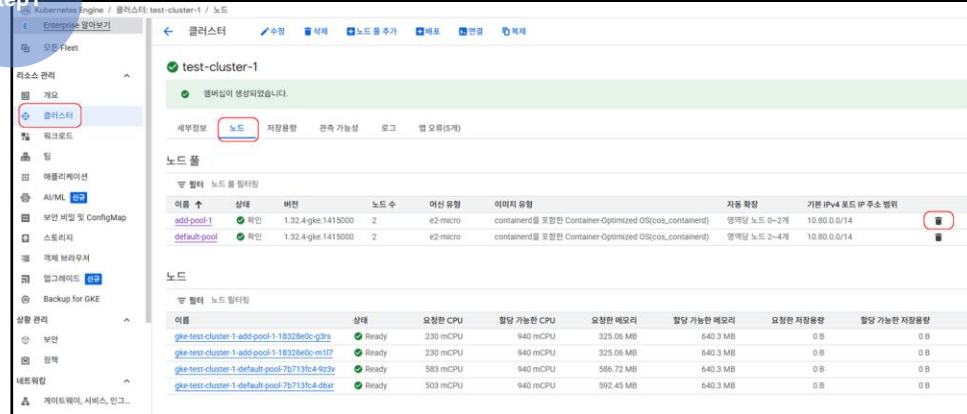


[그림 28] 노드 풀 수정

[노드 풀(Node Pool) 삭제]

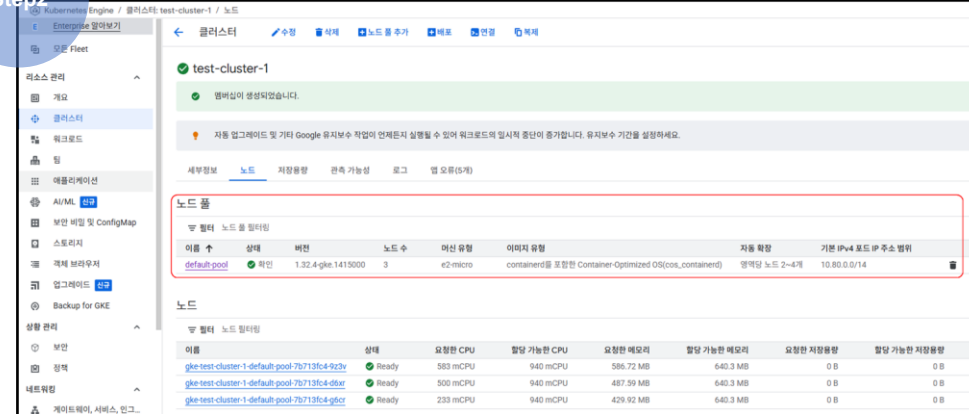
- 클러스터 > 대상 클러스터 이름 선택 > 노드(탭) > 삭제 대상 노트 풀(삭제)
- 노드 풀(리스트 확인)

Step1



[그림 29] 노드 풀 선택 및 삭제

Step2



[그림 30] 노드 풀 삭제 결과 확인



[AutoScaling-Pod]

- Horizontal Pod Autoscaler(HPA)

```
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: php-apache
spec:
  selector:
    matchLabels:
      run: php-apache
  replicas: 1
  template:
    metadata:
      labels:
        run: php-apache
    spec:
      containers:
      - name: php-apache
        image: k8s.gcr.io/hpa-example
        ports:
        - containerPort: 80
      resources:
        limits:
          cpu: 500m
        requests:
          cpu: 200m
```

[그림 31] deployment-php-apache.yml

```
---
apiVersion: autoscaling/v1
kind: HorizontalPodAutoscaler
metadata:
  name: php-apache
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: php-apache
  minReplicas: 1
  maxReplicas: 10
  targetCPUUtilizationPercentage: 50
```

[그림 32] hpa-php-apache.yml



[AutoScaling-Pod]

- Vertical Pod Autoscaler(VPA)

```
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: vpa-demo
spec:
  replicas: 1
  selector:
    matchLabels:
      app: vpa-demo
  template:
    metadata:
      labels:
        app: vpa-demo
    spec:
      containers:
        - name: stress-container
          image: vish/stress
          resources:
            requests:
              cpu: "50m"
              memory: "50Mi"
            limits:
              cpu: "200m"
              memory: "200Mi"
          args:
            - -cpus
            - "2"
```

[그림 33] vpa-demo-deployment.yaml

```
---
apiVersion: autoscaling.k8s.io/v1
kind: VerticalPodAutoscaler
metadata:
  name: vpa-demo
  namespace: default
  clusterName: gcp-edu-gke-standard-1
spec:
  targetRef:
    kind: Deployment
    name: vpa-demo
    apiVersion: apps/v1
  updatePolicy:
    updateMode: Auto
  resourcePolicy:
    containerPolicies:
      - containerName: "*"
        mode: Auto
        controlledResources:
          - cpu
          - memory
    minAllowed: {}
    maxAllowed: {}
```

[그림 34] vpa-deployment.yaml

```
---
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
metadata:
  name:
    verticalpodautoscalers.autoscaling.k8s.io
spec:
  group: autoscaling.k8s.io
  names:
    kind: VerticalPodAutoscaler
    listKind: VerticalPodAutoscalerList
    plural: verticalpodautoscalers
    singular: verticalpodautoscaler
    scope: Namespaced
  versions:
    - name: v1
      served: true
      storage: true
      schema:
        openAPIV3Schema:
          type: object
```

[그림 35] vpa-crd.yaml



[GCP Cloud Build 실습 목표]

- Cloud Build를 사용한 자동 빌드 & 배포
- Docker 이미지 자동 생성 및 GCR에 푸시
- GKE에 자동 배포 연결 (옵션)

[사전 준비]

- GCP 프로젝트 생성 및 Billing 활성화
- Cloud Build API 활성화
- GCR (Container Registry) 또는 Artifact Registry 활성화
- GitHub 또는 Cloud Source Repositories 연동



[예제 애플리케이션 구성]

- node.js 기반 간단한 웹 앱
- app.js 생성

```
const express = require('express');
const app = express();
const port = 8080;

app.get('/', (req, res) => res.send('Hello Cloud Build!'));

app.listen(port, () => console.log(`App listening on port ${port}`));
```

[그림 36] 노드 풀 삭제 결과 확인

[Dockerfile 생성]

```
FROM node:18

WORKDIR /app

COPY package*.json ./

RUN npm install

COPY . .

CMD ["node", "app.js"]
```

[그림 37] 노드 풀 삭제 결과 확인



[cloudbuild.yaml 작성]

```
steps:
- name: 'gcr.io/cloud-builders/docker'
  args: ['build', '-t', 'gcr.io/$PROJECT_ID/demo-app', '.']
- name: 'gcr.io/cloud-builders/docker'
  args: ['push', 'gcr.io/$PROJECT_ID/demo-app']

images:
- 'gcr.io/$PROJECT_ID/demo-app'
```

[그림 38] 노드 풀 삭제 결과 확인

[Cloud Build 트리거 생성]

- Cloud Console > Cloud Build > 트리거 생성
- GitHub 또는 Cloud Source Repo 연결
- Branch: main 또는 특정 브랜치
- 빌드 구성: cloudbuild.yaml 선택



[Cloud shell에 Github 연동]

- ssh key 생성

```
ssh-keygen -t rsa -b 4096 -C "your_email@example.com"
```

[그림 39] ssh key 생성

```
@cloudshell:~/ssh$ ssh-keygen -t rsa -b 4096 -C "jshan97@gmail.com"
Generating public/private rsa key pair.
Enter file in which to save the key (/home/jshan97/.ssh/id_rsa):
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /home/jshan97/.ssh/id_rsa
Your public key has been saved in /home/jshan97/.ssh/id_rsa.pub
The key fingerprint is:
SHA256:Sx8HceLwrf2gUP1b2qL0OfuYLeJyPFawtsBtpuSjRCU jshan97@gmail.com
The key's randomart image is:
+---[RSA 4096]-----+
|  . o . |
|  + *   |
| E * o   |
|  + +..  |
| S.o.=oo +|
| o ++*=.B |
|  oooB.o* |
| . = B .* |
| .. *.o+oo|
+-----[SHA256]-----+
```

[그림 40] ssh key 생성 완료



[Cloud shell에 Github 연동]

- ssh key 생성 공개키 확인

```
cat ~/.ssh/your_key_name.pub
```

[그림 41] ssh 공개키 확인

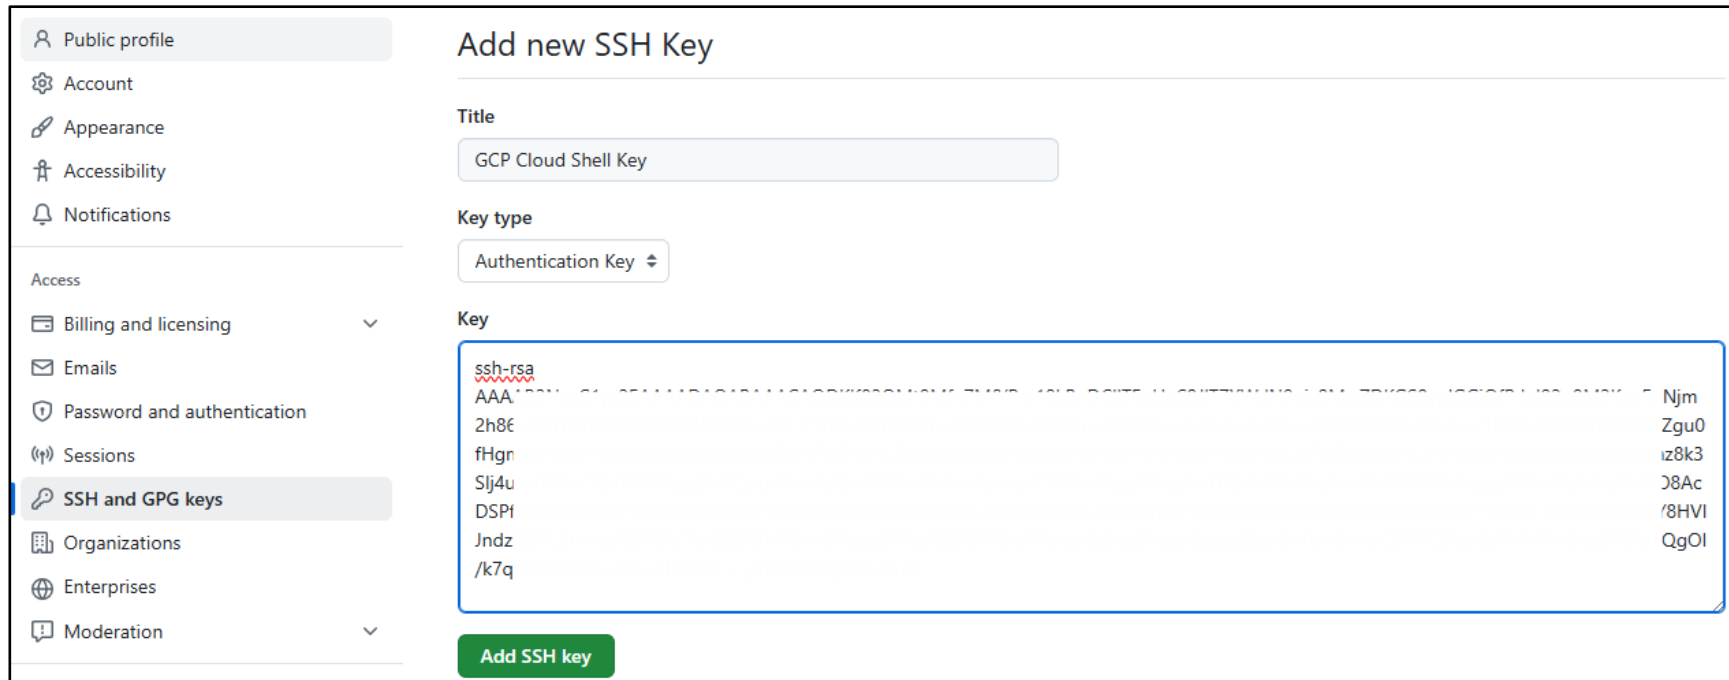
```
johanna@cloudshell:~/.ssh$ cat ~/.ssh/id_rsa.pub

ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAQCAQDK1f02OMt8Mfn7MfYPm18kBNDC11TEq
UpS8JlT7XWdN9niy8MwZDKGS8yp1GGiQfPdn182p8M2KpvFgNjm2h86xN/fITDrHBdlVe
C0I7/5eH6chJ+1fKhCdUb9WhN0WKZgu0fHgmQPF/0gPNIbpyMoBJaVeKriSSdZJV6FOzn
f/KuEqIioMs2AXaHoXFyDMH16UT91HckSHYhssMbQqFbfEw/s4xPhu/LH4zd/hx2VzmDw
kkggERcpyJmDQr4VfCh8FaZw+KHW94Melzggi9Wa69gURqIfUM08AcDSPfspeUiK34DvT
YzCWcxW0HS6P4cx5eyVzDLx6hNbNrE7ara2UAmqaWSwzUJhNjEBPFrGVc9QKSZnmVmi+j
PggojD73dGBDGW6V/79tApsTaBduBhopz8JpxuGSaHbNUwCtmaQPR6CDbpGUPVH9Ni64q
fXXDd1QgOI/k7qsdVMeEYh+AGp41XiPQ==
```

[그림 42] ssh 공개키 확인

[Cloud shell에 Github 연동]

- Github에 ssh-key 설정



Public profile

Account

Appearance

Accessibility

Notifications

Access

Billing and licensing

Emails

Password and authentication

Sessions

SSH and GPG keys

Organizations

Enterprises

Moderation

Add new SSH Key

Title

GCP Cloud Shell Key

Key type

Authentication Key

Key

```
ssh-rsa  
AAA... Njm  
2h8e Zgu0  
fHgr iz8k3  
Sij4u 38Ac  
DSPi /8HVI  
Jndz QgOI  
/k7q
```

Add SSH key

[그림 43] github ssh 공개키 등록



[Cloud shell에 Github 연동]

- Cloud shell ssh config 설정(~/.ssh/config)

```
Host github.com
  Hostname github.com
  IdentityFile ~/.ssh/your_key_name
```

```
Host github.com
  Hostname github.com
  IdentityFile ~/.ssh/id_rsa
```

[그림 44] ssh 사용 config 파일 설정

- Test the Connection

```
ssh -T git@github.com
```

```
[redacted]@cloudshell:~/.ssh$ ssh -T git@github.com
Hi [redacted]! You've successfully authenticated, but GitHub does not provide shell access.
```

[그림 45] ssh config 설정 확인



[Cloud shell에 git project repo 설정-1]

- git 초기화 / 사용자 이름 / 이메일주소 설정

```
git init
git config --global user.name "사용자이름"
git config --global user.email "이메일주소"
```

- 설정 확인

```
git config --list
```

```
j[redacted]@cloudshell:~/gcp-cloud-build-sample-1$ git config --list
core.repositoryformatversion=0
core.filemode=true
core.bare=false
core.logallrefupdates=true
credential.https://source.developers.google.com.helper=gcloud.sh
user.email=j[redacted]@[redacted]
user.name=j[redacted]
core.repositoryformatversion=0
core.filemode=true
core.bare=false
pull.rebase=false
remote.origin.url=git@github.com:gcp-edu/gcp-cloud-bulid-sample.git
remote.origin.fetch=+refs/heads/*:refs/remotes/origin/*
```

[그림 46] git 설정 확인



[Cloud shell에 git project repo 설정-2]

- git remote 설정

```
git remote add origin https://github.com/gcp-edu/gcp-cloud-bulid-sample.git
```

- 설정 확인


```
git remote -v
```

- Local branch 생성

```
git branch main
```

- 생성된 repo pull

```
git pull origin main
```



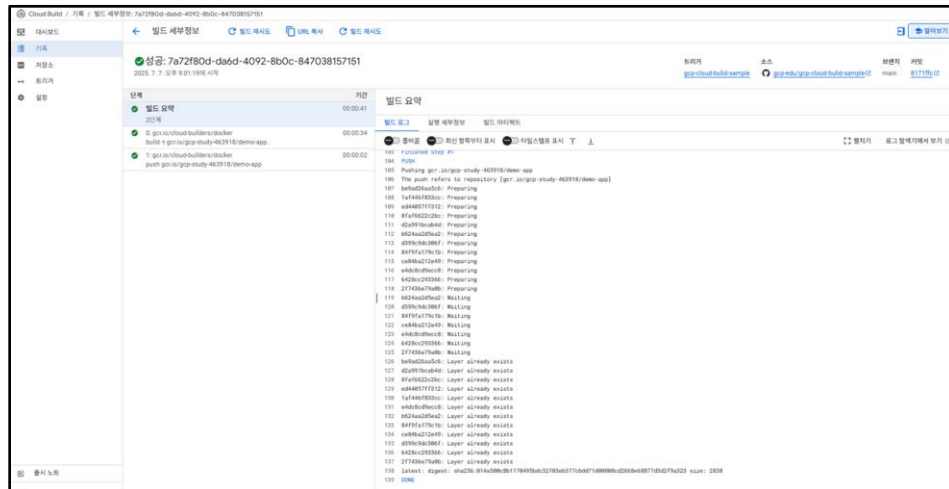
[그림 47] git repo ssh url 확인



[커밋 & 자동 실행 테스트]

- Cloud Build에서 자동으로 빌드 시작
- 이미지가 GCR로 푸시됨

```
git add .
git commit -m "Deploy with Cloud Build"
git push origin main
```



[그림 48] Cloud build 로그 확인



[그림 49] Artifact Registry 생성 확인



GKE에 Container 응용 배포하기

- GKE 클러스터 생성
- Docker 이미지 만들기
- Kubernetes YAML 작성
- GKE에 배포

사전 준비

- GCP 계정 필요
- 프로젝트 생성 (Billing enabled)
- Cloud Shell 또는 로컬 gcloud/kubectl 설치



1. GKE 클러스터 생성

- num-nodes=2
- region=us-central1
- machine-type=e2-medium
- disk-type=pd-standard
- disk-size=100

```
gcloud container clusters create demo-cluster \
  --num-nodes=2 \
  --region=us-central1 \
  --machine-type=e2-medium \
  --disk-type=pd-standard \
  --disk-size=100
```

```
$ gcloud container clusters create demo-cluster \
  --num-nodes=2 \
  --region=us-central1 \
  --machine-type=e2-medium \
  --disk-type=pd-standard \
  --disk-size=50
Note: The Kubelet readonly port (10255) is now deprecated. Please update
your workloads to use the recommended alternatives. See https://cloud.google.com/kubernetes-engine/docs/how-to/disable-kubelet-readonly-port for ways to check usage and for migration instructions.
Note: Your Pod address range ('--cluster-ipv4-cidr') can accommodate at most 1008 node(s).
Creating cluster demo-cluster in us-central1... Cluster is being health-checked (Kubernetes Control Plane is healthy)...done.

Created [https://container.googleapis.com/v1/projects/gcp-study-463918/zones/us-central1/clusters/demo-cluster].
To inspect the contents of your cluster, go to: https://console.cloud.google.com/kubernetes/workload_/gcloud/us-central1/demo-cluster?project=gcp-study-463918
kubeconfig entry generated for demo-cluster.
NAME: demo-cluster
LOCATION: us-central1
MASTER_VERSION: 1.32.4-gke.1415000
MASTER_IP: 34.67.148.58
MACHINE_TYPE: e2-medium
NODE_VERSION: 1.32.4-gke.1415000
NUM_NODES: 6
STATUS: RUNNING
```

[그림 50] demo-cluster 생성 command



2. app.js 파일 만들기

```
const express = require('express');
const app = express();
const port = 8080;

app.get('/', (req, res) => {
  res.send('Hello from GKE!');
});

app.listen(port, () => {
  console.log(`App is running on port ${port}`);
});
```

2. package.json 만들기

```
{
  "name": "demo-app",
  "version": "1.0.0",
  "main": "app.js",
  "scripts": {
    "start": "node app.js"
  },
  "dependencies": {
    "express": "^4.18.2"
  }
}
```



3. Dockerfile 만들기 & Push

- Dockerfile 작성

```
FROM node:18
WORKDIR /app
COPY package*.json ./
RUN npm install
COPY . .
CMD ["node", "app.js"]
```

- 생성 파일 구조

```
project-root/
├─ app.js
├─ package.json
└─ Dockerfile
```

- 빌드

```
#> docker build -t gcr.io/[PROJECT_ID]/demo-app:v1 .
```

- GCR Push

```
#> docker push gcr.io/[PROJECT_ID]/demo-app:v1
```

```
$ docker build -t gcr.io/gcp-study-463918/demo-app:v1 .
[+] Building 2.0s (10/10) FINISHED

    docker:default
=> [internal] load build definition from Dockerfile
0.0s
=> => transferring dockerfile: 134B
0.0s
=> [internal] load metadata for docker.io/library/node:18
0.6s
=> [internal] load .dockerignore
0.0s
=> => transferring context: 2B
0.0s
=> [1/5] FROM docker.io/library/node:18@sha256:c6ae79e38498325db67193d391e6eacd224d96c693a8a4d943498556716d3783
```

[그림 51] docker build command

```
$ docker push gcr.io/gcp-study-463918/demo-app:v1

The push refers to repository [gcr.io/gcp-study-463918/demo-app]
002762b6697c: Pushed
411422e32cb3: Pushed
4aab4bf90b47: Pushed
42d9c1bab012: Pushed
d2a991bcab4d: Layer already exists
b624aa2d5ea2: Layer already exists
d399c9dc306f: Layer already exists
84f9fa179c1b: Layer already exists
ce84ba212e49: Layer already exists
e4dc8cd9ecc8: Layer already exists
6428cc293366: Layer already exists
2f7436e79a0b: Layer already exists
v1: digest: sha256:4c2f3704c6d151a79688419e0267d3e54bec86b2016fec23d
f6387754d486e7 size: 2831
```

[그림 52] gcr push command



4. Kubernetes Deployment.yaml

- Deployment.yaml 작성

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: demo-app
spec:
  replicas: 2
  selector:
    matchLabels:
      app: demo-app
  template:
    metadata:
      labels:
        app: demo-app
    spec:
      containers:
        - name: demo-app
          image: gcr.io/[PROJECT_ID]/demo-app:v1
          ports:
            - containerPort: 8080
```

4. Service.yaml (LoadBalancer)

- Service.yaml 작성

```
apiVersion: v1
kind: Service
metadata:
  name: demo-service
spec:
  type: LoadBalancer
  selector:
    app: demo-app
  ports:
    - protocol: TCP
      port: 80
      targetPort: 8080
```



5. 적용 & 결과 확인

- `kubectl apply -f Deployment.yaml`

```
deployment.apps/demo-app created
```

- `Kubectl apply -f Service.yaml`

```
service/demo-service created
```

- `kubectl get pods`

NAME	READY	STATUS	RESTARTS	AGE
demo-app-764798b7d8-422vk	0/1	Error	4 (44s ago)	116s
demo-app-764798b7d8-zr9dp	0/1	CrashLoopBackOff	3 (32s ago)	116s

- `kubectl get svc`

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
demo-service	LoadBalancer	34.118.230.99	34.61.51.248	80:31572/TCP	2m25s
kubernetes	ClusterIP	34.118.224.1	<none>	443/TCP	32m



6. 환경 삭제

- `gcloud container clusters delete demo-cluster --region us-central1`

```
deployment.apps "demo-app" deleted
```

- `kubectl delete deployments demo-app`

```
$ gcloud container clusters delete demo-cluster --region us-central1
The following clusters will be deleted.
- [demo-cluster] in [us-central1]

Do you want to continue (Y/n)? y

Deleting cluster demo-cluster...done.
```



[Infrastructure as Code(IaC) 실습 목표]

- Terraform을 사용하여 VM 인스턴스를 배포
- IaC Tool : Terraform
- GCP Project 인증 정보 생성
 1. API 및 서비스 → 사용자 인증 정보 → 사용자 인증 정보 만들기 → 서비스 계정 → 역할 : 소유자
 2. 생성한 서비스 계정 클릭 → 키 → 키 추가 → 새 키 만들기 → 키 유형 : JSON → 인증키를 Terraform 폴더로 이동

The screenshot shows the Google Cloud Platform console for a service account named 'terraform-service-account'. The left pane displays the '서비스 계정 세부정보' (Service Account Details) tab, where the email address 'terraform-service-account@gcp-study-463918.iam.gserviceaccount.com' is highlighted with a red box. The right pane shows the '키' (Keys) tab, which contains a warning about key security and a table of existing keys. The table has columns for '유형' (Type), '상태' (Status), '키' (Key), '생성일' (Created), and '만료일' (Expires). One key is listed with a status of '활성' (Active) and its key ID is highlighted with a red box.

유형	상태	키	생성일	만료일
JSON	활성	73034e695420fdb41669a79d67bb0074dbccd206	2025. 7. 8.	10000. 1. 1.

[그림 53] terraform 서비스 계정 생성



Step 1. 인프라구조 정의(IaC Code 작성)

- main.tf 작성

```
# 1. Specify the Google Cloud Provider
# This block tells Terraform which cloud provider to use and how to
# authenticate.
# Ensure your gcloud CLI is authenticated and configured for the
# correct project.
terraform {
  required_providers {
    google = {
      source = "hashicorp/google"
      version = "~> 5.0" # Use a compatible version
    }
  }
}
```

[그림 54] main.tf (서비스 제공자 정의)

```
# 2. Configure the Google Cloud Project and Region/Zone
# You can define these as variables or hardcode them.
# Using variables makes your code more reusable.
variable "project_id" {
  description = "The Google Cloud Project ID."
  type        = string
  default     = "YOUR_PROJECT_ID" # <<< REMEMBER TO CHANGE THIS
}

variable "region" {
  description = "The Google Cloud region to deploy resources in."
  type        = string
  default     = "us-central1"
}

variable "zone" {
  description = "The Google Cloud zone to deploy the VM in."
  type        = string
  default     = "us-central1-c"
}

provider "google" {
  project = var.project_id
  region  = var.region
  credentials = "${file("????.json")}"
}
```

[그림 55] main.tf (인프라 구성 위치 및 대상 프로젝트 정보 설정)



Step 1. 인프라구조 정의(IaC Code 작성)

- main.tf 작성

```
# 3. Define the Compute Engine VM Instance
resource "google_compute_instance" "example_vm" {
  name     = "my-terraform-vm"
  machine_type = "e2-medium" # A cost-effective machine type
  zone     = var.zone

  boot_disk {
    initialize_params {
      image = "debian-cloud/debian-11" # Specify a public image
    }
  }

  network_interface {
    network = "default" # Use the default VPC network
    access_config {
      # This block creates an ephemeral external IP address
    }
  }
}
```

[그림 56] main.tf (VM 인스턴스 및 네트워크 정의)

```
# Optional: Add a startup script to install Nginx
metadata_startup_script = <<-EOF
#!/bin/bash
sudo apt-get update
sudo apt-get install -y nginx
echo "<h1>Hello from Terraform!</h1>" | sudo tee /var/www/html/index.nginx-debian.html
sudo systemctl start nginx
sudo systemctl enable nginx
EOF

tags = ["http-server"] # Apply tags for firewall rules
}
```

[그림 57] main.tf (install Nginx 정의)



Step 1. 인프라구조 정의(IaC Code 작성)

- main.tf 작성

```
# 4. Define a Firewall Rule to Allow HTTP Traffic
# This is crucial if you want to access the web server on the VM.
resource "google_compute_firewall" "allow_http" {
  name    = "allow-http-from-terraform"
  network = "default"

  allow {
    protocol = "tcp"
    ports    = ["80"]
  }

  source_ranges = ["0.0.0.0/0"] # Allow from anywhere (for demonstration)
  target_tags   = ["http-server"] # Apply to VMs with this tag
}
```

[그림 58] main.tf (방화벽 규칙 설정)

```
# 5. Output the VM's External IP Address
# This makes it easy to find the VM after deployment.
output "vm_external_ip" {
  value =
    google_compute_instance.example_vm.network_interface[0].access_config[0].nat_ip
  description = "The external IP address of the deployed VM."
}
```

[그림 59] main.tf (VM의 External IP 정의)



Step 2. 작업 디렉토리 초기화(terraform)

- Terraform 작업 디렉토리 초기화 명령

```
terraform init
```

```
gcp-study-463918@cloudshell:~/gcp-iac-terraform-sample (gcp-study-463918)$ terraform init

Initializing the backend...

Initializing provider plugins...
- Finding hashicorp/google versions matching "~> 5.0"...
- Installing hashicorp/google v5.45.2...
- Installed hashicorp/google v5.45.2 (signed by HashiCorp)

Terraform has created a lock file .terraform.lock.hcl to record the provider
selections it made above. Include this file in your version control repository
so that Terraform can guarantee to make the same selections by default when
you run "terraform init" in the future.

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

If you ever set or change modules or backend configuration for Terraform,
rerun this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.
```

[그림 60] terraform 초기화



Step 3. main.tf 적용 예정 내역 미리보기

terraform plan

```

$ terraform plan
Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:
+ create

Terraform will perform the following actions:

# google_compute_firewall.allow-http will be created
+ resource "google_compute_firewall" "allow-http" {
  + creation_timestamp = (known after apply)
  + destination_ranges = (known after apply)
  + direction          = (known after apply)
  + enable_logging     = (known after apply)
  + id                 = (known after apply)
  + name               = "allow-http-from-terraform"
  + network            = "default"
  + priority           = 1000
  + project            = "gcp-study-463918"
  + self_link          = (known after apply)
  + source_ranges      = [
    + "0.0.0.0/0",
  ]
  + target_tags        = [
    + "http-server",
  ]
  + allow {
    + ports = [
      + "80",
    ]
    + protocol = "tcp"
  }
}

# google_compute_instance.example-vm will be created
+ resource "google_compute_instance" "example-vm" {
  + can_ip_forward = false
  + cpu_platform   = (known after apply)
  + current_status = (known after apply)
  + deletion_protection = false
  + effective_labels = (known after apply)
  + guest_accelerator = (known after apply)
  + id             = (known after apply)
  + instance_id    = (known after apply)
  + label_fingerprint = (known after apply)
  + machine_type    = "e2-medium"
  + metadata_fingerprint = (known after apply)
  + metadata_startup_script = <<-EOF
    #!/bin/bash
    sudo apt-get update
    sudo apt-get install -y nginx
    echo "<h1>Hello from Terraform!</h1>" | sudo tee /var/www/html/index.nginx-debian.html
    sudo systemctl start nginx
    sudo systemctl enable nginx
  >>>EOF
}

```

[그림 61] plan 실행결과-1

```

    sudo systemctl enable nginx
  EOT
+ min_cpu_platform = (known after apply)
+ name             = "my-terraform-vm"
+ project          = "gcp-study-463918"
+ self_link        = (known after apply)
+ tags             = [
  + "http-server",
]
+ tags_fingerprint = (known after apply)
+ terraform_labels = (known after apply)
+ zone             = "asia-northeast3-b"

+ boot_disk {
  + auto_delete = true
  + device_name = (known after apply)
  + disk_encryption_key_sha256 = (known after apply)
  + kms_key_self_link = (known after apply)
  + mode            = "READ_WRITE"
  + source          = (known after apply)

  + initialize_params {
    + image = "debian-cloud/debian-11"
    + labels = (known after apply)
    + provisioned_iops = (known after apply)
    + provisioned_throughput = (known after apply)
    + size = (known after apply)
    + type = (known after apply)
  }
}

+ network_interface {
  + internal_ipv6_prefix_length = (known after apply)
  + ipv6_access_type           = (known after apply)
  + ipv6_address               = (known after apply)
  + name                       = (known after apply)
  + network                   = "default"
  + network_ip                 = (known after apply)
  + stack_type                 = (known after apply)
  + subnetwork                 = (known after apply)
  + subnetwork_project         = (known after apply)

  + access_config {
    + nat_ip = (known after apply)
    + network_tier = (known after apply)
  }
}

Plan: 2 to add, 0 to change, 0 to destroy.

Changes to Outputs:
+ vm_external_ip = (known after apply)

```

[그림 62] plan 실행결과-2



Step 4. main.tf 배포

```
terraform apply
```

```
Changes to Outputs:
+ vm_external_ip = (known after apply)

Do you want to perform these actions?
Terraform will perform the actions described above.
Only 'yes' will be accepted to approve.

Enter a value: yes

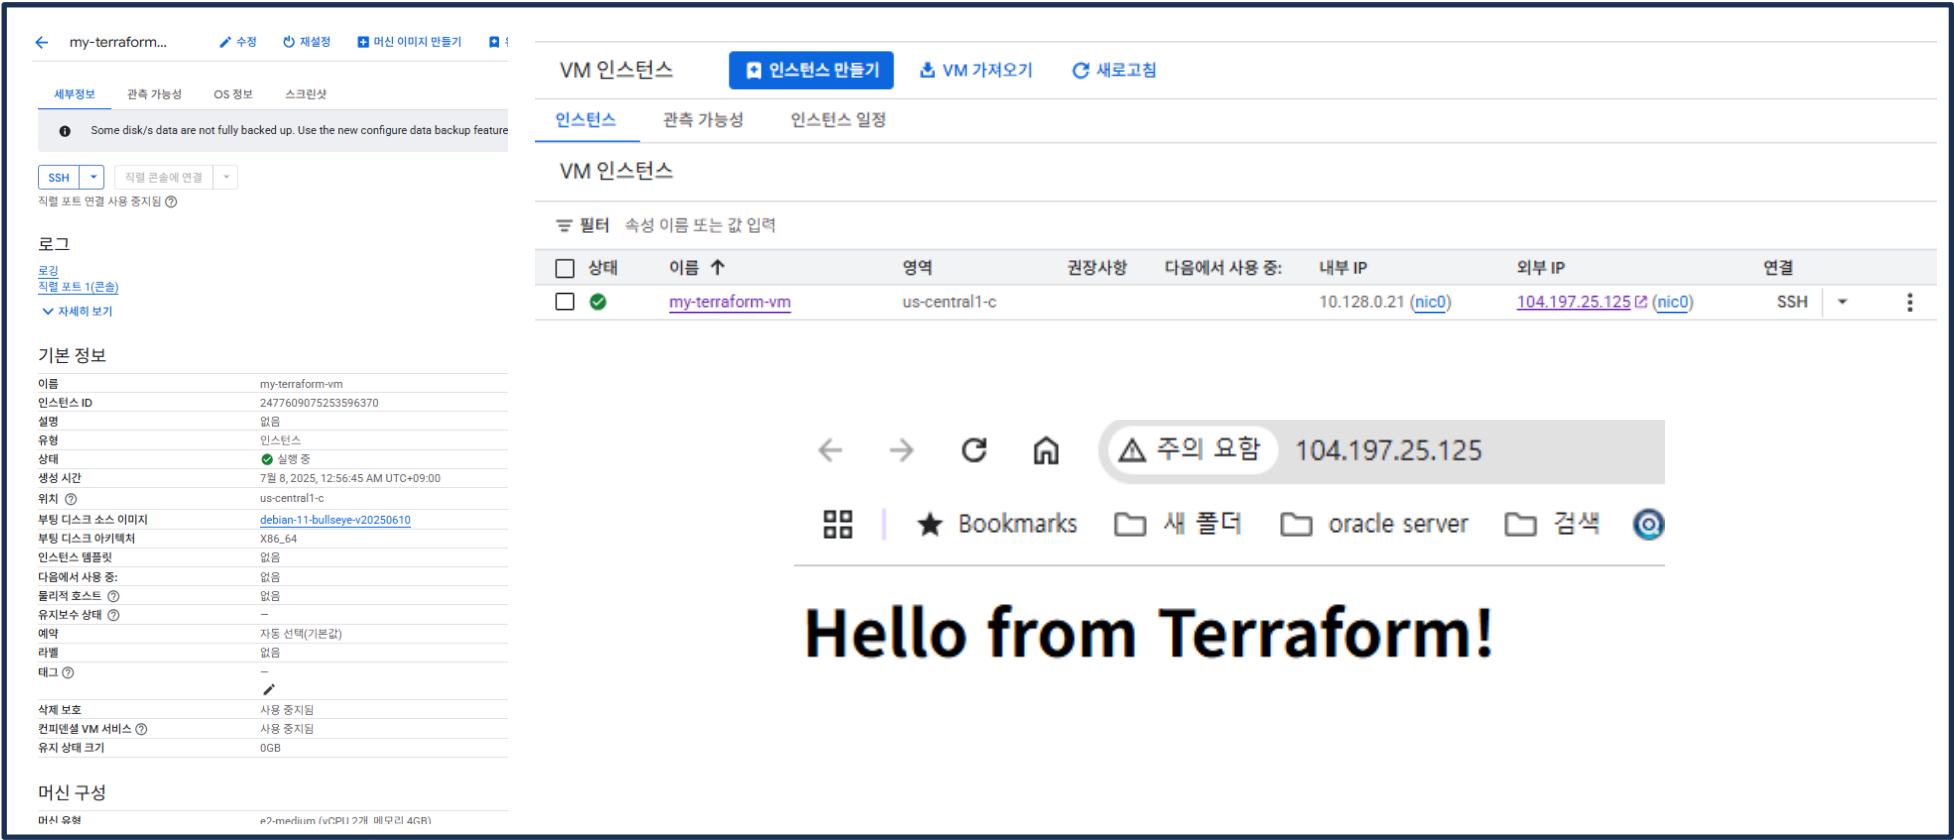
google_compute_firewall.allow_http: Creating...
google_compute_instance.example_vm: Creating...
google_compute_firewall.allow_http: Still creating... [10s elapsed]
google_compute_instance.example_vm: Still creating... [10s elapsed]
google_compute_firewall.allow_http: Creation complete after 12s [id=projects/gcp-study-463918/global/firewalls/allow-http-from-terraform]
google_compute_instance.example_vm: Still creating... [20s elapsed]
google_compute_instance.example_vm: Creation complete after 28s [id=projects/gcp-study-463918/zones/us-central1-c/instances/my-terraform-vm]

Apply complete! Resources: 2 added, 0 changed, 0 destroyed.

Outputs:
vm_external_ip = "104.197.25.125"
```

[그림 63] apply 실행결과

Step 4. main.tf 배포 결과



[그림 64] apply 배포결과 확인



Step 5. main.tf 구성 삭제

terraform destroy

```
- scheduling {
-   automatic_restart = true -> null
-   min_node_cpus     = 0     -> null
-   on_host_maintenance = "MIGRATE" -> null
-   preemptible       = false -> null
-   provisioning_model = "STANDARD" -> null
- }
- shielded_instance_config {
-   enable_integrity_monitoring = true -> null
-   enable_secure_boot         = false -> null
-   enable_vtpm                = true  -> null
- }
- }

Plan: 0 to add, 0 to change, 2 to destroy.

Changes to Outputs:
  vm_external_ip = "104.197.25.125" -> null

Do you really want to destroy all resources?
  Terraform will destroy all your managed infrastructure, as shown above.
  There is no undo. Only 'yes' will be accepted to confirm.

Enter a value: yes

google_compute_firewall.allow_http: Destroying... [id=projects/gcp-study-463918/global/firewalls/allow-http-from-terraform]
google_compute_instance.example_vm: Destroying... [id=projects/gcp-study-463918/zones/us-central1-c/instances/my-terraform-vm]
google_compute_firewall.allow_http: Still destroying... [id=projects/gcp-study-463918/global/firewalls/allow-http-from-terraform, 10s elapsed]
google_compute_instance.example_vm: Still destroying... [id=projects/gcp-study-463918/zones/us-central1-c/instances/my-terraform-vm, 10s elapsed]
google_compute_firewall.allow_http: Destruction complete after 12s
google_compute_instance.example_vm: Still destroying... [id=projects/gcp-study-463918/zones/us-central1-c/instances/my-terraform-vm, 20s elapsed]
google_compute_instance.example_vm: Still destroying... [id=projects/gcp-study-463918/zones/us-central1-c/instances/my-terraform-vm, 30s elapsed]
google_compute_instance.example_vm: Still destroying... [id=projects/gcp-study-463918/zones/us-central1-c/instances/my-terraform-vm, 40s elapsed]
google_compute_instance.example_vm: Still destroying... [id=projects/gcp-study-463918/zones/us-central1-c/instances/my-terraform-vm, 50s elapsed]
google_compute_instance.example_vm: Still destroying... [id=projects/gcp-study-463918/zones/us-central1-c/instances/my-terraform-vm, 1m0s elapsed]
google_compute_instance.example_vm: Still destroying... [id=projects/gcp-study-463918/zones/us-central1-c/instances/my-terraform-vm, 1m10s elapsed]
google_compute_instance.example_vm: Still destroying... [id=projects/gcp-study-463918/zones/us-central1-c/instances/my-terraform-vm, 1m20s elapsed]
google_compute_instance.example_vm: Still destroying... [id=projects/gcp-study-463918/zones/us-central1-c/instances/my-terraform-vm, 1m30s elapsed]
google_compute_instance.example_vm: Still destroying... [id=projects/gcp-study-463918/zones/us-central1-c/instances/my-terraform-vm, 1m40s elapsed]
google_compute_instance.example_vm: Still destroying... [id=projects/gcp-study-463918/zones/us-central1-c/instances/my-terraform-vm, 1m50s elapsed]
google_compute_instance.example_vm: Still destroying... [id=projects/gcp-study-463918/zones/us-central1-c/instances/my-terraform-vm, 2m0s elapsed]
google_compute_instance.example_vm: Still destroying... [id=projects/gcp-study-463918/zones/us-central1-c/instances/my-terraform-vm, 2m10s elapsed]
google_compute_instance.example_vm: Still destroying... [id=projects/gcp-study-463918/zones/us-central1-c/instances/my-terraform-vm, 2m20s elapsed]
google_compute_instance.example_vm: Destruction complete after 2m24s

Destroy complete! Resources: 2 destroyed.
```

[그림 65] destroy 결과 확인